

Report

DQN

First attempt standard DQN with the set of params.
The algorithm is taken from the DQN exercise and then modified.

The basic implementation follow the original paper: <https://arxiv.org/pdf/1312.5602.pdf>

The model used consists of a neural network with two hidden layers respectively with 64 and 32 units. It is bigger than what is necessary and the reason is that a bigger network helps with the problem of moving targets.

To help with this problem the target network is also implemented.

To help with the problem of having experience samples not identically distributed a simply replay buffer is implemented.

I used the RMSProp optimiser with LR 0.0005.

Exploration strategy EpsilonGreedy with epsilon start of 1.0 and final 0.01 with a decay of 0.995


Other parameters (standard worked well):

```
BUFFER_SIZE = int(1e5)    # replay buffer size
BATCH_SIZE = 64           # minibatch size
GAMMA = 0.99              # discount factor
TAU = 1e-3                # for soft update of target parameters
LR = 5e-4                 # learning rate
UPDATE_EVERY = 4          # how often to update the network
```

I train the agent for 1000 episodes and each episode made of 500 time steps

With this setup the agent is able to complete the task with an average reward over 100 episodes greater than 13.0 in 494 episodes.

```
Episode 100    Average Score: 0.79
Episode 200    Average Score: 5.10
Episode 300    Average Score: 8.32
Episode 400    Average Score: 10.01
Episode 494    Average Score: 13.01
Environment solved in 494 episodes!    Average Score: 13.01
Episode 500    Average Score: 13.27
Episode 600    Average Score: 14.71
Episode 700    Average Score: 15.21
Episode 800    Average Score: 16.06
Episode 900    Average Score: 16.24
Episode 1000   Average Score: 15.52
```



Improvements:

1. Double DQN

The first improvement was to use a double DQN following the paper: <https://arxiv.org/pdf/1509.06461.pdf>

This helps in mitigating the overestimation of action - value function of DQN decoupling action selection from action evaluation.

2. Dueling architecture

Another improvement is shown in the paper: <https://arxiv.org/pdf/1511.06581.pdf>

This improvement uses the fact that actions of a state are correlated and when we learn from one we can use the information to learn something about the other actions of the same state.

3. Prioritized Experience Replay

We use experience replay because we need to have identically distributed samples across experience.

But we also know some samples provide more information than others as explained in this paper:

<https://arxiv.org/pdf/1511.05952.pdf>

Note: the implementation is inspired from the book: *Grokking Deep Reinforcement Learning* (<https://www.manning.com/books/grokking-deep-learning>)

Parameters (standard from the book):

$\alpha=0.6$,

$\beta_0=0.4$,


$\beta_{\text{rate}}=0.99992$

Epsilon: 0.5 # used for weighted importance-sampling

We train for the same amount of time and episodes as before.

The agent is able to win in 477 episodes. The learning is more stable and the score higher.

Episode 100	Average Score: 0.51
Episode 200	Average Score: 3.85
Episode 300	Average Score: 6.92
Episode 400	Average Score: 9.79
Episode 477	Average Score: 13.07
Environment solved in 477 episodes! Average Score: 13.07	
Episode 500	Average Score: 13.21
Episode 600	Average Score: 14.37
Episode 700	Average Score: 15.23
Episode 800	Average Score: 16.50
Episode 900	Average Score: 16.40
Episode 1000	Average Score: 16.87



Further improvements:

There are several improvements possible and i plan to work on them later.

1. Noisy DQN: <https://arxiv.org/pdf/1706.10295.pdf>
2. Distributional Q-Learning: <https://arxiv.org/pdf/1707.06887.pdf>
3. Asynchronous Learning: <https://arxiv.org/pdf/1602.01783.pdf>

We can also improve the Replay Buffer using the Segment Tree as shown in the OpenAI baseline:

<https://github.com/openai/baselines/tree/master/baselines>

DQN from raw pixels:

I provided the same implementation of the improved DQN to work with visual observations.

The main difference are:

1. Observations reshaped from (1,3,84,84) to (3,84,84)
2. Stacked frames input(12,84,84)
3. Convolutional Neural Network

CNN architecture:

The input is passed through 3 convolution layers with 32, 64, 64 filter.

After the convolutions there is a fully connected layer and from there the standard Dueling Architecture is followed.

There are several improvements that could be done.

As example skipping frames.

Unfortunately i have no GPUs and i was not able to try this version. It also doesn't work with the Udacity Workspace. I tried using Google Colab but so far it did not work. I decided to move to Atari Games.